



Rohit Aggarwal <rohit313agg@gmail.com>

AI Agents Have Hit A New Gear

AI Breakfast <aibreakfast@mail.beehiiv.com>

Fri, Jul 4, 2025 at 9:54 AM

Reply-To: AI Breakfast <aibreakfastemail@gmail.com>

To: "rohit313agg@gmail.com" <rohit313agg@gmail.com>

July 04, 2025 | [Read Online](#)

AI BREAKFAST

AI Agents Have Hit A New Gear

Good morning. It's Friday, July 4th.

On this day in tech history: In **1969**, the Stanford Research Institute developed [Shakey the Robot](#), which used a combination of computer vision, a problem-solving program called STRIPS, and basic AI to navigate and interact with its environment, pioneering autonomous robotics despite its limited 12-bit computer processing.

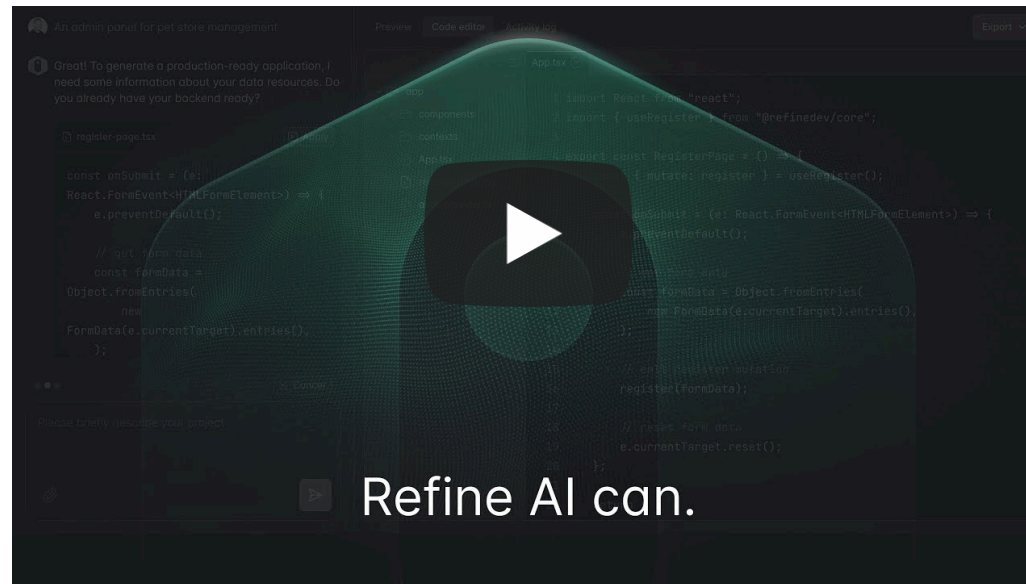
In today's email:

- OpenAI Denies Token, Gets Major Oracle Power
- AI Agents Have Hit A New Gear
- Meta Poaches SSI CEO After Failed Buyout
- 5 New AI Tools
- Latest AI Research Papers

You read. We listen. Let us know what you think by replying to this email.

In partnership with React AI

Build Internal Enterprise Apps at Warp Speed with AI



RefineAI can.

Unlock the Future of Enterprise App Development.

No-code is history. React from scratch? That's madness. Vibe coding is the future—and generic AI builders often fall short. Unless it's [RefineAI](#).

Forget black-box outputs and zero maintainability. [RefineAI](#) uses domain knowledge and the open-source Refine framework (31K Stars) to generate production-grade internal React apps.

- ✅ **Go beyond simple prompts!** Extract complex project specifications from API docs and more.
- ✅ **Vibe coding your internal software.** Use natural language to add any CRUD functionality, new resources and components.
- ✅ **Download or deploy!** Download your project to continue development or deploy it with a single click.

✅ **Choose your own UI!** With the Ant design, MUI and shadcn/ui options and more to come.

Try it now!

Thank you for supporting our sponsors!

AI NEWS

Today's trending AI news stories

OpenAI Locks 4.5 GW Oracle Buildout, Denies Robinhood Token Ties

OpenAI is scaling compute and tightening control over its brand and data pipelines. It has secured [4.5 GW of dedicated capacity from Oracle](#) to power its Stargate AI systems, part of a \$30B/year cloud commitment. Oracle will expand hyperscale infrastructure in Texas, Michigan, and Wyoming, doubling the Abilene site's footprint to 2 GW. The deployment ranks among the largest AI-dedicated compute expansions globally, underscoring the hardware demands of next-gen foundation models.



Stargate data center at Abilene, Texas | Image: w.media

Concurrently, OpenAI disavowed “OpenAI tokens” launched on Robinhood, calling them [unaffiliated and unauthorized](#). The tokens, issued via an SPV (special purpose vehicle), falsely imply equity access. OpenAI clarified no shares were sold or transferable without consent.



OpenAI Newsroom
@OpenAINewsroom



[These “OpenAI tokens” are not OpenAI equity. We did not partner with Robinhood, were not involved in this, and do not endorse it. Any transfer of OpenAI equity requires our approval—we did not approve any transfer.](#)

Please be careful.

8:06 PM • Jul 2, 2025

13.3K Likes 867 Retweets

652 Replies

In contrast, OpenAI is backing the [ChatGPT Data Collective](#) with [Vana](#), where users upload chat logs, earn \$GPT tokens, and vote on data usage, signaling support for on-chain data governance amidst intensifying AI–web traffic tensions. [Read more](#).

AI Agents hit a new gear, running end-to-end and full stack with no humans needed

A new wave of agentic tools dropped this week, replacing manual workflows with autonomous AI execution.

HeyGen’s [Video Agent](#) handles full video production from a sentence or file, writing scripts, generating footage, casting virtual actors, and editing with pro-level timing. Touted as the first true “Creative Operating System,” it targets ad creators, product teams, and content marketers looking for speed at scale.



Joshua Xu
@joshua_xu_



Today we're introducing you to the future of video.

The world's first Creative Operating System, we call it the HeyGen Video Agent.

Upload a doc, some footage, or even just a sentence.

It analyzes your input.

Finds the story.

Writes the script with taste.

Selects the shots



3:05 PM • Jun 26, 2025

13K Likes 1.37K Retweets

585 Replies

Replit's upgraded Dev Agent, '[Dynamic Intelligence](#),' adds long-range reasoning, high-context understanding, live web search, and multi-step autonomy, writing APIs, debugging, and refactoring with minimal input.



 [Introducing Dynamic Intelligence for Replit Agent](#)

We're unlocking next-level problem solving with:

 [Extended Thinking](#)

 [High Power Model](#)

 [Web Search](#)

Now the Agent thinks deeper, reasons better, and surfs the web to get you complete solutions.

This allows Agent



5:26 PM • Jul 2, 2025

546 Likes 67 Retweets

47 Replies

Genspark has launched [AI Docs](#), completing its trio of agentic tools alongside AI Slides and Sheets. It automatically generates rich, markdown-ready documents from simple prompts, handling structure, formatting, and flow with no templates or manual input.



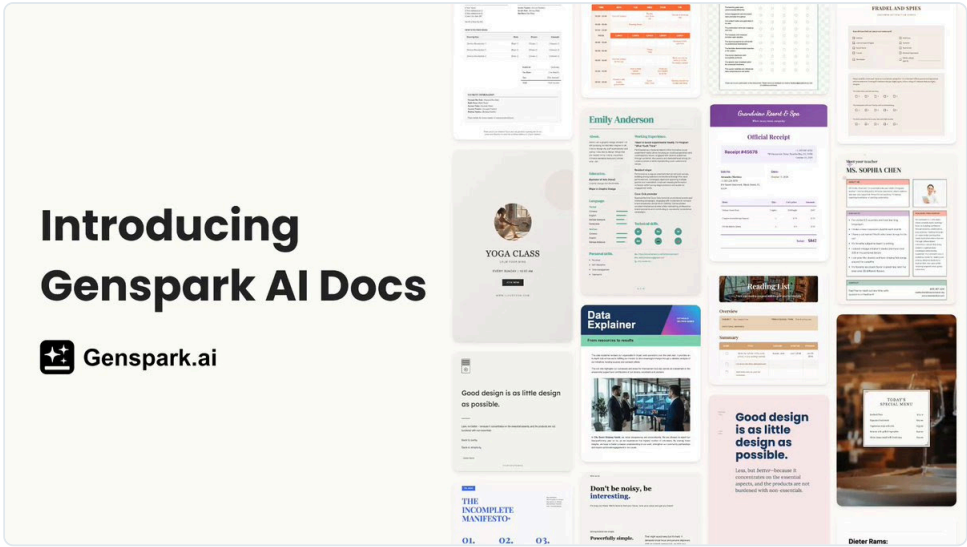
Genspark
@genspark_ai



 **BIG NEWS:** Introducing Genspark AI Docs — the world's first full-agentic AI document creator with native support for both rich text AND markdown!

With AI Slides, AI Sheets, and now AI Docs, our classic productivity trio for the AI era is finally complete! 🎉

Simply tell



12:50 PM • Jul 2, 2025

234 Likes 34 Retweets

12 Replies

And [Shortcut](#), from Fundamental, tackles complex Excel workflows, solving championship-grade tasks 10x faster than human pros. In trials, it solved Excel World Championship Cases with over 80% accuracy in under 10 minutes, operating with zero formulas or macros.

**nico**

@nicochristie



Introducing Shortcut — the first superhuman Excel agent.

Shortcut one-shots most knowledge work tasks on Excel.

It even scores >80% on Excel World Championship Cases in ~10 minutes. That's 10x faster than humans.

Our early preview is live. Just comment for an invite code.



4:00 PM • Jul 2, 2025

7.37K Likes **574** Retweets

[3.78K Replies](#)

Meta and SSI Lock Horns in Superintelligence Race

Superintelligence Inc. (SSI) just got a major shake-up. Co-founder [Daniel Gross](#) has exited the company and is now leading Meta's new superintelligence division, Meta Superintelligence Labs (MSL), alongside Nat Friedman and ex-Scale AI CEO Alexander Wang. This comes after Meta's failed attempt to buy SSI outright for a reported \$32 billion.



Daniel Gross / Nat Friedman | Image: The Information

With Gross out, former OpenAI chief scientist Ilya Sutskever now runs SSI as CEO while continuing to lead its technical direction. Sutskever reiterated SSI’s singular mission: to build safe superintelligence without dilution from commercial product demands. Still, Gross’s departure raises questions about internal dynamics and timing, especially if SSI is nearing a major technical inflection point.



Ilya Sutskever
@ilyasut



I sent the following message to our team and investors:

As you know, Daniel Gross’s time with us has been winding down, and as of June 29 he is officially no longer a part of SSI. We are grateful for his early contributions to the company and wish him well in his next

3:58 PM • Jul 3, 2025

11.8K Likes 646 Retweets

601 Replies

Now, with both Meta and SSI racing toward superintelligence, one with a centralized commercial engine, the other with a singular research mandate, the competition between their labs frames a defining chapter in the AI industry’s next wave. [Read more.](#)

ETCETERA

Stories you may have missed

- [Adding calendar events with a screenshot is AI at its finest](#)
- [Black Forest Labs opens its AI image model FLUX.1 context \[dev\] for private use](#)
- [Meta Eyes Trillions With Chatbots That Engage Unprompted](#)
- [Sakana AI Introduces AB-MCTS: Letting AIs Think as a Team](#)
- [People are using AI to 'sit' with them while they trip on psychedelics](#)
- [Grok Will Soon Let You Create Games](#)
- [A project to bring CUDA to non-Nvidia GPUs is making major progress](#)
- [Perplexity AI launches new 'Max' subscription priced at \\$200/month](#)
- [Red Teams Testing "Claude-Neptune-v3": Claude 4.5 in Sight?](#)
- [Watch: Longevity Tech CEO: 120-Year Lifespans in 20 Years, Escape Velocity in 50](#)
- [Nvidia briefly on track to become world's most valuable company ever](#)
- [NYT to start searching deleted ChatGPT logs after beating OpenAI in court](#)
- [Google rolls out its new Veo 3 video-generation model globally](#)

- AI virtual personality YouTubers, or 'VTubers,' are earning millions
- Microsoft laying off about 9,000 employees in latest round of cuts
- Baidu launches AI video generator, overhauls search features
- This AI model was trained on 10M human choices. Now it thinks and reacts like us
- US engineers' new way of attacking vision systems can make AI see whatever you want
- No one likes meetings. They're sending their AI note takers instead


AI TOOLS

5 new AI-powered tools from around the web

[create-api.dev](#)

A lightweight AI-powered tool for generating OpenAPI specs. No login, no setup. Built by Kong to make spec creation and collaboration faster with AI. It's a first step, with more planned around how AI can improve the full API lifecycle.

[create-api.dev](#)



LLM Gateway

A fully open-source AI Gateway for your LLM needs, route, manage, and analyze your LLM requests across multiple providers with a unified API interface.

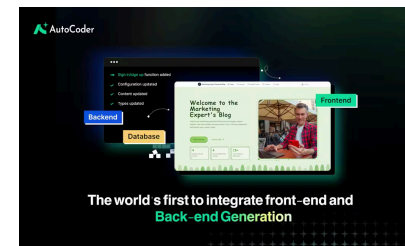
llmgateway.io



AutoCoder

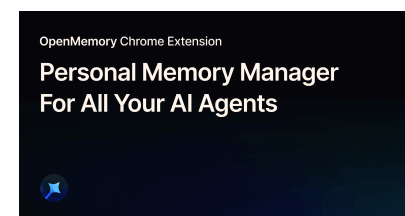
The 1st full stack vibe coding tool. Ask #AutoCoder to build a website = Frontend + backend + Database. No need for Supabase.

www.autocoder.cc/platform



OpenMemory Chrome Extension

OpenMemory is a personal AI memory manager that allows you to manage your preferences across ChatGPT, Claude, Perplexity, and




[more.](#)

[openmemory.app](#)

[a0.dev](#)

World's fastest iOS builds + one-click App Store submission! AI agent with thinking mode, payments integration, Stripe support, Github sync & more. Turn your ideas into real apps in seconds.



[a0.dev](#)

AI RESEARCH

[arXiv](#) is a free online library where researchers share pre-publication papers.

 [Does Math Reasoning Improve General LLM Capabilities? Understanding Transferability of LLM Reasoning](#)


 [Kwai Keye-VL Technical Report](#)

 [LongAnimation: Long Animation Generation with Dynamic Global-Local Memory](#)

 [LangScene-X: Reconstruct Generalizable 3D Language-Embedded Scenes with TriMap Video Diffusion](#)

 [Skywork-Reward-V2: Scaling Preference Data Curation via Human-AI Synergy](#)





Thank you for reading today's edition.



Your feedback is valuable. Respond to this email and tell us how you think we could add more value to this newsletter.

Interested in reaching smart readers like you? To become an AI Breakfast sponsor, reply to this email or DM us on X!



Update your email preferences or unsubscribe [here](#)

© 2025 AI Breakfast

228 Park Ave S, #29976, New York, New York 10003, United States

 Powered by beehiiv

